

Insurance Claim Processing Hub

Team BrainWave, KL University — ServiceNow Hackathon (Insurance Track)

What it is

An insurance claims app built on ServiceNow. A customer submits a claim, AI reads and checks the documents, a fraud score and settlement suggestion are generated automatically, and a human adjuster (with a manager for riskier cases) makes every real decision. The AI never decides anything on its own — it only prepares information for a person to act on.

The problem

Claims today are handled manually: documents are checked by hand, there's no consistent way to spot fraud, settlement amounts vary adjuster to adjuster, and customers have to call in just to ask what's happening. This app fixes each of those.

How it works

1. **Customer submits a claim** through a portal form — own policy only, enforced automatically.
2. **AI reads the documents** (Gemini), checks nothing's missing, and flags anything that doesn't match up (e.g. a date on the police report that doesn't match the claim).
3. **Fraud score** is calculated from clear, listed reasons (a Decision Table maps the score to Low/Medium/High).
4. **Settlement range** is suggested by comparing the claim to similar past claims.
5. **An adjuster** reviews everything, can ask the customer for more info, and recommends approve or reject with an amount or reason.
6. **A manager** signs off only when the risk is high or the amount is well outside the suggested range — otherwise this step is skipped.
7. **The adjuster finalizes** — writes the real outcome message to the customer, and the claim closes.

Throughout, the customer only ever sees plain status updates — never the fraud score, AI notes, or internal comments.

What's built

Area	Built with
Data & app structure	App Engine
Customer portal & cases	Customer Service Management (CSM)
Fraud scoring	Decision Table + Business Rules
Document reading & summaries	Now Assist / Gemini AI
Reporting	Native Reports & Dashboard

Not used, despite being listed as available tools in the brief: FSO Flow Designer, IntegrationHub, and ServiceNow's built-in Approvals Engine. Their jobs were done instead with Business Rules, a direct AI API call, and a custom sign-off flow. This was a deliberate choice to fit the build within the time available, not an oversight.

Testing

About 93 test cases were run across the whole app — document checks, AI reading, fraud scoring, settlement suggestions, the full adjuster/manager workflow, and reporting. All passed in the final run. Eleven real bugs were found and fixed along the way (for example: a rule that silently did nothing, a note that wasn't actually reaching the customer, and a currency symbol showing "\$" instead of "Rs."). Full detail is in the separate Test Plan.

What's not done yet

- If a customer replies to an adjuster's request, the claim doesn't automatically move back into review — the adjuster has to notice.
- Emails are generated correctly but don't actually send (no mail server set up on this instance).
- Three of the decision buttons use a pop-up input box that once failed unpredictably elsewhere in the app; they've tested fine but haven't been rebuilt with the safer pattern.
- The claim form still uses ServiceNow's default layout, not a custom-designed one.

Bottom line

The core idea — AI prepares the evidence, a human always decides — works end to end and is tested. The technology list wasn't followed exactly; the app was built with a smaller, working set of tools instead, and that trade-off is stated here plainly rather than hidden.