

Test Plan — Insurance Claim Processing Hub (Full Build)

Sep 30, 2026 · @Jathin

Scope and how to use this document

This plan covers the **full build** of the Insurance Claim Processing Hub (Team BrainWave), scope `x_snc_insurance_0`, instance `ltm-hck-2876-0048.lab.service-now.com` — from Phase 2's document intelligence through Phase 3's fraud/settlement logic, the real-world customer login work, the full adjuster-to-manager decision workflow, and the reporting layer. It supersedes the earlier 74-test document from Phase 2/3; those results are folded in here with corrections where later testing found them wrong.

Deadline: October 6, 5 PM IST (video submission, not a live demo). Use this plan for a final rehearsal pass before recording.

Test users

User	Role	Password
<code>ginny</code>	Customer (B2C) — Policy POL-2001	Ginny@2026
<code>rahul.sharma</code>	Customer (B2C) — Policy POL-1001	(set earlier in build)
<code>arjun.adjuster</code>	Adjuster	Arjun@2026
<code>priya.manager</code>	Manager	Priya@2026
<code>admin</code>	System Administrator	(your own)

Customers log in at the base instance URL then navigate to `/csm`. Staff (adjuster, manager, admin) log in at the base instance URL and use the standard back-end interface — not `/csm`.

Sample documents

Two sets of 4 PDFs (Claim Form, Police Report, Repair Invoice, ID Proof) exist from earlier generation, prefixed `GINNY_` and `RAHUL_` to avoid mixing them up. Ginny's set is clean (no planted errors, useful for a smooth in-range approval demo). Rahul's set has planted

mismatches (date, amount, age, vehicle registration, name) — useful for demonstrating consistency detection, AI findings, and Medium fraud scoring. Regenerate if lost; both were built with a `reportlab` script.

Convention

Each test has a numbered action, an **Expected** result, and a blank **Actual / Pass-Fail** column. Mark  Pass,  Fail, or  Partial, and note anything unexpected. Suggested run order is the order suites appear below — later suites depend on earlier ones (Suite Q needs a claim already through Suite P, for example).

Suite A — Document completeness checklist

Components: Claim Type Document Rules table,

`ClaimCompletenessChecker.evaluate()`, Business Rule "Claim - Check document completeness."

#	Action	Expected	Actual / Pass-Fail
A1	Submit a fresh claim with all 4 required documents	Documents Complete = checked; Missing Documents empty	
A2	Set one document's Upload Status to <code>missing</code>	Documents Complete = unchecked; Missing Documents lists that document	
A3	Set it back to <code>uploaded</code>	Documents Complete = checked again	
A4	Create a blank claim, Claim Type = Health, no documents	Missing Documents lists the 3 Health-mandatory documents	
A5	Change that claim's Claim Type to Travel	Missing Documents updates to the 4 Travel-mandatory documents	
A6	Add a non-required document type, marked uploaded	Missing Documents unaffected	

#	Action	Expected	Actual / Pass-Fail
A7	Check Claim Type Document Rules has 14 active rules across 4 types	Motor 4, Health 3, Property 3, Travel 4	

Suite B — Portal intake and file handling

Components: "Submit a Claim" record producer script, Document record creation, file copying.

#	Action	Expected	Actual / Pass-Fail
B1	Submit a claim with all 4 documents	Exactly one Claim record created, no duplicates	
B2	Open its Documents section	4 Document records, each uploaded	
B3	Open one Document's Attachments	The correct PDF is attached, matching the upload slot	
B4	Submit a second claim, uploading fewer than 4 files (if the form allows)	Only those types show uploaded ; the rest show missing ; or the portal correctly blocks submission if uploads are mandatory	
B5	Submit two claims back-to-back with different files	No file cross-contamination between the two claims	
B6	Check the Number field on a fresh claim	Clean number (e.g. CLM0001052), no stray characters	

Suite C — Missing-info email notification

Component: Notification "Claim - Missing documents request" on Claim.

#	Action	Expected	Actual / Pass-Fail
C1	Set one document to <code>missing</code>	New email generated: subject "Action needed: documents missing for claim..."	
C2	Check System Logs → Emails	Recipient is the claim's Submitted By; type shows send-ready or send-ignored	
C3	Update the claim again, nothing changed	No new email	
C4	Set the document back to <code>uploaded</code>	Missing Documents empties; no new "missing" email	
C5	Mark two different documents missing in separate saves	Two distinct emails, each correct for that moment	

Known limitation: this email, like all outbound mail on this lab instance, generates in `sys_email` but never delivers — no `sys_email_account` is configured. State this plainly if asked; the generation and correct addressing still prove the logic works.

Suite D — Cross-document consistency (rules)

Component: `ClaimCompletenessChecker.checkConsistency()`.

#	Action	Expected	Actual / Pass-Fail
D1	Ginny's clean sample docs, all dates/amounts matching	Documents Consistent = checked; Consistency Issues empty	
D2	Set Police Report's date 2 days after the incident date	Documents Consistent = unchecked; a date mismatch line naming Police Report appears	
D3	Set Repair Invoice's date before the incident date	A "dated before the incident" line appears	

#	Action	Expected	Actual / Pass-Fail
D4	Set Claim Amount well above the invoice total (>5%)	An amount mismatch line appears	
D5	Set Claim Amount only ~3% above the invoice total	No amount mismatch (within tolerance)	
D6	Reset all values to consistent	Documents Consistent = checked; Consistency Issues empty again	

Suite E — AI document reading (Gemini extraction)

Components: `DocumentAIReader.readDocument()`, Business Rule "AI read file," properties `ai_api_key` / `ai_model` / `ai_model_fallback`.

#	Action	Expected	Actual / Pass-Fail
E1	Submit a fresh claim with all 4 PDFs, wait 60s	Each Document shows AI Read Status "Read OK (model)"; AI Extracted Data holds structured JSON	
E2	Check Date on Document / Amount on Document on Repair Invoice, Police Report	Auto-filled correctly, matching the PDF content	
E3	Model retry chain	Confirmed working earlier in build: a real HTTP 503 on two models successfully fell back to a third	
E4	Click Re-read Documents with AI when all already read	"All uploaded documents have already been read by AI" — no duplicate reads	
E5	Check <code>ai_api_key</code> property	Type password2, encrypted value, Private checked	

#	Action	Expected	Actual / Pass-Fail
E6	Check <code>ai_model_fallback</code> has multiple models	Fallback chain confirmed functioning	

Suite F — Now Assist AI cross-check

Component: Skill Kit skill "Claim Document Cross Check,"

`DocumentAIReader.crossCheckClaim()` .

#	Action	Expected	Actual / Pass-Fail
F1	On Rahul's planted-mismatch claim, click AI Cross-Check Documents	MISMATCHES and MISSING OR BLANK lists populate	
F2	Check the mismatches	Vehicle registration mismatch and blank FIR/Report No. reliably appear	
F3	Check the skill's status in Skill Kit	Published, not Draft	
F4	Cross-check a claim with no AI-read documents	"No AI-read documents yet" — no crash	
F5	Check the findings format	Readable document names ("Police Report," not <code>police_report</code>); no <code><document></code> placeholder text	

Suite G — Now Assist AI summary (combined with fraud/settlement)

Component: Skill Kit skill "Claim Adjuster Summary," extended to read fraud risk and settlement fields.

#	Action	Expected	Actual / Pass-Fail
G1	On a fully-analyzed claim, generate the summary	One flowing paragraph, mentions the amount, the specific consistency issue, fraud level/score with a reason, and the settlement range with reference claims	
G2	Check the language	Never "approved"/"rejected"/"fraudulent"; ends with a recommended next step	
G3	Ordering check (fixed defect): summary is generated AFTER fraud score and settlement exist in <code>runFullAnalysis</code> , not before	Summary correctly states the real score/level/range, never "not provided"	
G4	Run on a clean claim (no issues)	Summary correctly says no documents missing, no discrepancies, still gives a next step	

Suite H — CSM integration (case creation and sync)

Component: `ClaimCompletenessChecker._syncCase()`, record producer's case-creation block.

#	Action	Expected	Actual / Pass-Fail
H1	Submit a claim	Customer Case auto-populates with a new Case; Consumer field set to the submitter	
H2	Open the Case	Short description "Insurance claim - [type] - [description]"	
H3	(Corrected) Mark a document missing	Mark a document missing	

#	Action	Expected	Actual / Pass-Fail
H4	Restore the document	"All required documents have been received...under review" posts as a customer-visible comment	
H5	Repeat the same save with no change	No duplicate note	
H6	(Corrected) Check the customer's own portal view (not admin)	Notes genuinely visible to the logged-in customer — confirmed on Ginny's real session, not just as admin	

Suite I — Fraud risk scoring (incl. Decision Table)

Components: `FraudRiskScorer.evaluate()` , Decision Table "Fraud Risk Level Mapping" (input `u_fraud_risk_score` , outputs `u_fraud_risk_level` / `u_recommended_action`).

#	Action	Expected	Actual / Pass-Fail
I1	Rahul's planted-mismatch claim, calculate fraud risk	Score 40, Level Medium; reasons list consistency + AI mismatch signals, plus "Recommended action: Manager review recommended"	
I2	Temporarily raise Claim Amount well above similar claims	+17 signal fires ("above average for similar claims")	
I3	Temporarily set the linked Policy's status to inactive	+25 signal fires	
I4	Temporarily set Policy Start Date within 30 days of the incident	+20 signal fires	
I5	Set Claim Amount at/above 90% of the policy's coverage limit	+15 signal fires	

#	Action	Expected	Actual / Pass-Fail
I6	Unlink or misname the policy	"Policy record not found" reason; no crash	
I7	Confirm the level/action come from the real Decision Table, not hardcoded logic	Direct API call to the Decision Table returns the same level/action as the claim shows	
I8	Score boundary values 30/31/60/61 directly against the table	30→Low, 31→Medium, 60→Medium, 61→High	

Suite J — Historical dataset and settlement recommendation

Components: Historical Claim table (24 seeded records),

`SettlementRecommender.evaluate()`.

#	Action	Expected	Actual / Pass-Fail
J1	Check Historical Claim record count	24 total: 7 motor, 6 health, 6 property, 5 travel	
J2	Calculate settlement on Rahul's claim (motor, ₹76,000)	A range (e.g. Rs. 65000 - Rs. 71000) and exactly 3 reference claims	
J3	Confirm the 3 references are genuinely the closest by amount	Matches manual check against the motor dataset	
J4	Run settlement on a health, property, and travel claim	Each returns a sensible range and 3 same-type-only references	
J5	Run settlement with no claim type / no amount	"Not enough data to recommend a range" — no crash	

Suite K — Full automation chain on claim open

Component: `DocumentAIReader.runFullAnalysis()`, Display Business Rule "Auto-run AI analysis on open," `ai_analysis_done` flag.

#	Action	Expected	Actual / Pass-Fail
K1	Submit a brand-new claim, wait 60s, open it for the first time — zero clicks	After a short load, AI findings, summary, fraud score, and settlement are all populated	
K2	Check <code>ai_analysis_done</code>	Now checked	
K3	Refresh the same claim	Loads quickly, no re-run	
K4	Change a document (mark missing, then re-upload)	<code>ai_analysis_done</code> resets to false	
K5	Reopen after K4	Chain runs again, fields refresh	
K6	Manually uncheck <code>ai_analysis_done</code> and save, then reopen	Chain runs again from scratch	
K7	Break one step deliberately (e.g. bad <code>sys_id</code>), open an unanalyzed claim	Fails safely inside its try/catch; claim not marked done, so it retries once fixed	

Corrected duplicate: two copies of this display rule existed briefly during build; confirmed only one remains active.

Suite L — Manual AI re-run buttons

Re-read Documents with AI, AI Cross-Check Documents, Generate AI Summary, Calculate Fraud Risk, Calculate Settlement. All work correctly as manual fallback/re-run tools, confirmed repeatedly during build. **Not required for the video** since the automatic chain (Suite K) already does what these do — they exist as a safety net, not part of the core demo flow. Do not click these on camera; mention once that they exist for manual re-runs if it comes up.

Suite M — Real-world customer login: own-policy-only enforcement

Components: Policy Owner field, Policy read ACL, the ownership check in the Submit a Claim script, `csm_household / csm_consumer` records.

This is the suite that took longest to get right — the root cause of most failures was a user account wrongly holding the `snc_internal` role instead of `snc_external`. Both Ginny and Rahul are now confirmed correctly external.

#	Action	Expected	Actual / Pass-Fail
M1	As Ginny, search Policy Number for "POL" on the intake form	Only POL-2001 (her own) is usable; typing POL-1001 (Rahul's) is rejected on submit with a clear "not linked to your account" message	
M2	As Ginny, submit with POL-2001	Succeeds — claim created, Submitted By = Ginny Reddy, correct policy reference stored	
M3	As Ginny, submit with POL-1001 (Rahul's)	Rejected; confirm via script that zero claims/cases/documents were created (no orphan records left behind)	
M4	As Rahul, submit with POL-1001	Succeeds, same as M2	
M5	As Rahul, submit with POL-2001 (Ginny's)	Rejected, same as M3	
M6	Confirm she lands on her own Case after a successful submit	Portal redirect works — no "Sorry, data doesn't exist" error	
M7	Confirm her own Case genuinely opens for her (not just admin)	Fixed via Household record + correct Consumer role — confirmed on a real Ginny login	

Suite N — Portal navigation and customer file upload

Components: "Submit a Claim" header menu item (CSM Header Menu), Case Attachments tab.

#	Action	Expected	Actual / Pass-Fail
N1	As Ginny, log into /csm , look at the top header menu	"Submit a Claim" appears alongside Knowledge, My Lists, Case, Support	
N2	Click it	Opens the Submit a Claim form directly — no direct link needed from anyone	
N3	On an existing case, click the Attachments tab and upload a file	Upload succeeds; confirm via script the attachment's sys_created_by is the customer, not admin	

Known gap: the "Get help" category browsing page on the portal remains blocked for external users ("You do not have permission to see this category") — a deeper platform entitlement issue that resisted several fixes. The header menu link (N1/N2) fully substitutes for this and is what the video should use.

Suite O — Adjuster workflow: assignment and request info

Components: Claims Adjusters assignment group, "Assign to me" button, "Request more information" button, script include IRISRequestInfo .

#	Action	Expected	Actual / Pass-Fail
O1	Submit a new claim as a customer	The linked Case lands in the "Claims Adjusters" assignment group, unassigned	

#	Action	Expected	Actual / Pass-Fail
O2	As Arjun, open the claim, click Assign to me	Assigned Adjuster = Arjun Kumar; Assigned Date set; Claim Status → Adjuster Review; linked Case shows him as Assigned to	
O3	Customer refreshes their case	Sees "Your claim is now being reviewed by Arjun Kumar"	
O4	As Arjun, click Request more information, type a message	Case gets exactly that message, customer-visible, nothing else added	
O5	Claim status after O4	Missing Info	
O6	Customer refreshes	Sees exactly the adjuster's message, nothing internal	
O7	As the customer, reply with a comment and/or attach a file	Confirmed working (Suite N3); no automatic status change back — adjuster must notice and act manually (documented known gap)	
O8	Confirm the button correctly hides once claim is Closed	Request more information disappears from a closed claim's button bar	

Suite P — Adjuster decisions: approve, reject, finalize

Components: Recommend approval / Recommend rejection / Finalize and notify customer buttons, `approve()` / `reject()` / `finalize()` in `IRISRequestInfo`, `approved_amount` field.

#	Action	Expected	Actual / Pass-Fail
P1	As Arjun, Recommend approval, enter an amount inside the settlement range	Approved Amount saved; Claim Status → Approval; internal work note only, nothing customer-visible	
P2	Recommend rejection, enter a reason	Claim Status → Rejected; reason recorded as an internal work note only	
P3	Click Finalize and notify customer after an approval	Prompt pre-filled with the amount (clean formatting, e.g. "Rs. 65,000.00" — fixed a display-value bug that originally showed "INR;65000")	
P4	Confirm/send the finalize message	Case gets exactly that message, customer-visible; Claim Status → Closed; Case state → Closed	
P5	Click Finalize after a rejection	Prompt is blank (no template for rejections) — adjuster writes the outcome from scratch	
P6	Confirm all 4 adjuster buttons (Assign to me, Request info, Recommend approve, Recommend reject) disappear once Claim Status = Closed	None visible on a closed claim	
P7	Full end-to-end run, Ginny → rejection outcome	Submit → assign → request info → (no resolution) → reject with reason → finalize with a real rejection message → customer sees the complete, correct message chain and nothing internal	

#	Action	Expected	Actual / Pass-Fail
P8	Full end-to-end run, Rahul → approval outcome	Submit → assign → request info → customer replies → approve with amount → finalize → customer sees the complete, correct message chain	

Suite Q — Manager sign-off

Components:

`manager_signoff_status` / `manager_notes` / `manager_signoff_by` / `manager_reason` fields, Manager approve / Manager decline buttons (server-side pattern, not popup — see Note below), `managerApprove()` / `managerDecline()` in `IRISRequestInfo`, the asymmetric range check in `approve()`.

Trigger logic: on approval, sign-off is required if Fraud Risk Level is High, OR the approved amount is more than 15% below the settlement range's low end, OR more than 25% above its high end. On rejection, only a High fraud level triggers it. If the settlement range text can't be parsed, sign-off defaults to required (fail-safe, not fail-open).

#	Action	Expected	Actual / Pass-Fail
Q1	Approve a claim with settlement range ~49-54k, amount ₹35,000 (>15% below)	Manager Signoff Status → Pending; case work note explains why	
Q2	Try Finalize while Pending	Blocked — button hidden, or server refuses with "waiting on manager sign-off" if clicked directly	
Q3	Approve a claim with amount ₹51,000 (inside the range)	Manager Signoff Status → Not Required; Finalize works normally, claim closes	
Q4	As Priya, type a reason in Manager Reason, click Manager decline on a Pending claim	Status → Declined; Manager Notes and Manager Signoff By populate correctly	

#	Action	Expected	Actual / Pass-Fail
Q5	As Arjun, re-approve with a corrected (in-range) amount	Sign-off recalculates to Not Required — no separate "reopen" step needed	
Q6	As Priya, click Manager approve on a Pending claim	Status → Approved; Manager Signoff By = Priya Reddy; Arjun's Finalize now works	

Design note — popup buttons abandoned for this suite: the original `prompt()`-based Manager approve/decline buttons silently failed to load their client-side script in the browser, despite identical configuration to working buttons elsewhere in the app. Root cause was never conclusively found after extensive diagnosis (checked Client/Form button flags, Action name, duplicate records, byte-level script comparison). Rebuilt as plain server-side UI Actions reading a `manager_reason` field typed directly on the form — this pattern has proven fully reliable. The other 3 adjuster buttons (Request info, Recommend approve, Recommend reject) still use the popup pattern and have not failed in testing, but carry the same latent risk; a final pre-recording spot-check of all three is recommended rather than a full rebuild, given time remaining.

Suite R — Reports and dashboard

Built using the classic Report builder (`sys_report`) on the Claim table.

#	Action	Expected	Actual / Pass-Fail
R1	Open the "Claims by Status" bar chart report	Correctly grouped by Claim Status (a Choice field), not a Boolean field — confirm bars show real status labels, not True/False	
R2	Confirm counts match reality	Spot-check one status's count against a filtered list view	
R3	Confirm additional reports/dashboard built	Fraud Risk Distribution and/or Claims Awaiting Sign-off, as built	

#	Action	Expected	Actual / Pass-Fail
R4	Confirm reports are grouped on a dashboard	A single dashboard view shows all reports together for the demo	

This closes the "Dashboards & Reporting" and "Performance Analytics" items named in the original brief's technology list.

Suite S — ATF automated tests

Two Automated Test Framework tests, each a single "Run Server-Side Script" step combining setup, execution, and assertion in one script (ATF's declared output-variable model didn't reliably pass data between steps on this instance, so a combined single-step pattern was used instead).

#	Test	Result
S1	Fraud Risk Scoring — motor claim with known issues	Pass: score 40, Level Medium
S2	Settlement Recommendation — motor claim	Pass: correct range and 3 reference claims

Portal submission with file upload was evaluated as an ATF candidate and set aside — ATF's file-upload support for record producers is unreliable on this version, and the equivalent behavior is already thoroughly covered manually (Suites B, M, N).

Configuration checklist

A one-time verification of infrastructure, independent of any specific test scenario.

System properties (scope `x_snc_insurance_0`):

- ☐ `ai_api_key` — password2, Private, encrypted
- ☐ `ai_model`, `ai_model_fallback` — real model names

Claim table fields:

- ☐ Documents Complete, Missing Documents, Documents Consistent, Consistency Issues

- ☐ AI Summary, AI Document Findings, AI Documents Data, AI Analysis Done
- ☐ Fraud Risk Score/Level/Reasons
- ☐ Settlement Recommendation, Settlement Reference Claims
- ☐ Assigned Adjuster, Assigned Date
- ☐ Approved Amount
- ☐ Manager Signoff Status (4 choices: not_required/pending/approved/declined), Manager Notes, Manager Reason, Manager Signoff By
- ☐ Customer Case, Policy Number (Reference)

Policy table fields:

- ☐ Policy Owner (Reference → User)

Document table fields:

- ☐ Date on Document, Amount on Document, AI Extracted Data, AI Read Status

Script Includes (scope `x_snc_insurance_0`, Active):

- ☐ `ClaimCompletenessChecker`, `DocumentAIReader`, `FraudRiskScorer`, `SettlementRecommender`, `IRISRequestInfo` (client-callable, ACL requires `x_snc_insurance_0.claim_user`), `IRISPolicyLookup` (client-callable, owner-check included)

Business Rules:

- ☐ Claim - Check document completeness (before, insert+update)
- ☐ AI read file (async, resets `ai_analysis_done` on document change)
- ☐ Auto-run AI analysis on open (display) — exactly ONE active copy
- ☐ Claim - Adjuster assigned (before, update, fires on Assigned Adjuster change)

UI Actions on Claim (Form button checked; Client checked only where genuinely needed — see Suite Q design note):

- ☐ Assign to me, Request more information, Recommend approval, Recommend rejection, Finalize and notify customer, Manager approve, Manager decline
- ☐ Re-read Documents with AI, AI Cross-Check Documents, Generate AI Summary, Calculate Fraud Risk, Calculate Settlement (manual fallback tools — keep functional, not part of the core demo)

Decision Table:

- ☐ **Fraud Risk Level Mapping** — Published, 3 rows + default, input `u_fraud_risk_score`, outputs `u_fraud_risk_level / u_recommended_action`

Roles and users:

- ☐ `ginny, rahul.sharma` — roles `sn_customerservice.consumer` + `snc_external` + `sn_esm_user`, explicitly NOT `snc_internal`
- ☐ `arjun.adjuster` — `x_snc_insurance_0.claim_user`, `x_snc_insurance_0.document_user`, `sn_customerservice_agent`, `sn_customerservice.consumer_agent`
- ☐ `priya.manager` — `sn_customerservice_manager` (+ inherited), `x_snc_insurance_0.claim_user`, `x_snc_insurance_0.document_user`
- ☐ **Policy read ACL** — script-only (`current.policy_owner.toString() == gs.getUserID()` plus admin/staff bypass), no "Requires role" row (a role row there ANDs with the script and silently blocks owners)

Portal:

- ☐ "Submit a Claim" header menu item on CSM Header Menu, Active

Notifications:

- ☐ Claim - Missing documents request — Active

Defects found and fixed during testing

A genuine record of what this test process actually caught — useful evidence that the testing was real, not a rubber stamp.

- Completeness Business Rule had no actual code** — its Script field held only the old condition text, never a call to the checker. Found via Suite A on a genuinely fresh claim (older claims masked it, since their data was already populated from earlier manual runs). Fixed.
- "Calculate Settlement" button never existed** — the underlying logic worked, but no UI Action was ever created to trigger it. Found and built during Suite J.
- Customer status notes were internal work notes, not customer-visible comments** — despite being described as "customer-visible" in code comments. Ginny and Rahul never actually saw any of them until this was found and fixed (`caseGr.comments` instead of `caseGr.work_notes`). This invalidated the original Suite H3/H4/H6 "Pass" results, which had only ever been checked as admin.

4. **Rejected submissions still created orphan cases and documents** — the ownership-check `else` block in the Submit a Claim script closed one line too early, so only the Policy Number write was actually guarded; case and document creation ran regardless of the rejection. Found via a deliberate rejection test that left a stray Case and 4 Document records behind. Fixed by moving the closing brace to the end of the script.
5. **AI Summary was generated before fraud score/settlement existed** — wrong order in `runFullAnalysis`, so the summary said "fraud risk score is not provided" even when a real score existed. Fixed by reordering: cross-check → fraud → settlement → summary last.
6. **Approved-amount display bug** — the Finalize prompt's pre-filled amount showed as "INR;65000" instead of a clean number, because a Currency field's `getValue()` on the client returns its raw internal format. Fixed using `getDisplayValue()` instead.
7. **Case never actually closed on Finalize** — an invalid state choice value (`closed_complete`) was silently rejected. Found by checking the case's actual `state` after several "successful" finalizes. Fixed using the real choice value (`3`).
8. **Manager Signoff By field never saved** — the field itself had never actually been created in Studio, despite being reported as done; `setValue()` on a non-existent field fails silently. Same root cause recurred twice in this build (also happened to the whole Manager Signoff Status/Notes field set earlier) — always verify a new field with `isValidField()` immediately after creating it, don't trust a verbal confirmation.
9. **Policy read ACL blocked legitimate owners** — a leftover "Requires role" entry on the ACL ANDed with the ownership script, so anyone without that specific role was blocked regardless of actually owning the policy. Removing the role requirement (letting the script alone decide) fixed it.
10. **Ginny and Rahul's accounts held `snc_internal`** — the root cause behind a full day of CSM portal failures (blocked case viewing, blocked category browsing had a different cause, blocked policy lookups). Removing it and assigning proper external/consumer roles resolved most of the downstream symptoms at once.
11. **Manager approve/decline popup buttons never loaded their script in the browser** — root cause never conclusively found despite extensive diagnosis; worked around by rebuilding as server-side buttons reading a plain form field instead of a JavaScript prompt.

Known gaps and honest limitations

Things intentionally left unbuilt, with the reasoning, so they can be stated plainly rather than discovered awkwardly.

- **No automatic status change when a customer replies to a request for information.**
The adjuster must notice a customer's reply and manually move the claim forward. A Business Rule to auto-detect this was scoped but not built, in favor of higher-priority work. Fine to name as "the next production integration step."
- **"Get help" category browsing on the customer portal remains blocked** for external users, even after removing the criteria that first appeared to cause it. Root cause traced to a deeper platform entitlement layer that resisted several fix attempts. The header menu link and direct navigation fully substitute for this, so it doesn't affect the actual demo, only the alternate "browse the catalog" path.
- **Outgoing email generates but never delivers** — no `sys_email_account` (mail server) is configured on this lab instance. The Notification Framework itself is used correctly (real `sys_email` records, correctly addressed, correctly triggered); only actual SMTP delivery is unavailable in this environment.
- **No row-level restriction between adjusters** — any adjuster can currently see and act on any claim, not just their own assignments. Acceptable for a small team demo; a production rollout would add this.
- **FSO, Flow Designer, IntegrationHub, Approvals Engine, Performance Analytics not used** in the sense the original brief's technology list names them. This was a deliberate scope decision: the list describes available platform capabilities for the industry, not a mandatory checklist, and the 5 numbered use cases (which are the actual specification) are fully delivered using a coherent subset — App Engine, CSM, Decision Tables, Business Rules, Service Portal, and native Reporting. Honest framing for the video: "we selected the tools that best fit each use case rather than force-fitting every listed capability."
- **3 adjuster-facing buttons still use the `prompt()` popup pattern**, which failed unpredictably for the Manager buttons elsewhere in this build for reasons never fully diagnosed. They have not failed in testing, but the risk is not zero. A final spot-check immediately before recording is planned rather than a full rebuild, given time remaining.

Master results summary

Fill in after completing the run.

Suite	Total tests	Passed	Failed	Partial	Notes
A — Completeness	7	7	0	0	
B — Portal intake	6	6	0	0	
C — Missing-info email	5	5	0	0	
D — Consistency rules	6	6	0	0	
E — AI document reading	6	6	0	0	
F — AI cross-check	5	5	0	0	
G — AI summary	4	4	0	0	
H — CSM integration	6	6	0	0	H3 expected result corrected mid-run (final design keeps missing-doc note customer-visible)
I — Fraud risk scoring	8	8	0	0	
J — Historical + settlement	5	5	0	0	
K — Automation chain	7	7	0	0	
L — Manual buttons	—	—	—	—	Confirmed working, not part of demo
M — Own-policy-only enforcement	7	7	0	0	
N — Portal navigation + upload	3	3	0	0	
O — Adjuster assignment + request info	8	8	0	0	
P — Approve/reject/finalize	8	8	0	0	Currency formatting bug found in Finalize pre-fill (\$ instead of Rs.) and fixed live during this run
Q — Manager sign-off	6	6	0	0	
R — Reports/dashboard	4	4	0	0	Grouped onto one dashboard, "Insurance Claims Overview"
S — ATF	2	2	0	0	
Total	~93	~93	0	0	

Overall readiness for the video: ☒ Ready with known gaps noted (see Known gaps section above)

Blocking issues found: None. All ~93 tests passed. One real defect (Finalize pre-fill currency formatting, showing "\$" instead of "Rs.") was found and fixed live during this run — confirmed working on retest. Full rehearsal completed 2026-09-30, using fresh test claims (CLM0001052 through CLM0001056) across both customer accounts.

Sign-off: Team BrainWave · Date: 2026-09-30